



Lecture 4: Artificial Neural Networks: Multilayer Networks and BP

Xuming He
SIST, ShanghaiTech
Fall, 2019


Outline

- Multi-layer neural networks
 - Limitations of single layer networks
 - Networks with single hidden layer
 - Sequential network architecture
- Inference and learning
 - Forward and Backpropagation
 - Examples: one-layer network
 - General BP algorithm

Capacity of single neuron

■ Binary classification

- A neuron estimates $P(y = 1|\mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x})$
- Its decision boundary is linear, determined by its weights



Capacity of single neuron


- Can solve linearly separable problems

$$\mathcal{D} = \mathcal{D}^+ \cup \mathcal{D}^-$$

$$\exists \mathbf{w}^*, \mathbf{w}^{*\top} \mathbf{x} > 0, \forall \mathbf{x} \in \mathcal{D}^+$$


$$\mathbf{w}^{*\top} \mathbf{x} < 0, \forall \mathbf{x} \in \mathcal{D}^-$$

- Examples



Capacity of single neuron


- Can't solve non linearly separable problems



- Can we use multiple neurons to achieve this?


Capacity of single neuron

- Can't solve non linearly separable problems
- Unless the input is transformed in a better representation



Capacity of single neuron

- Can't solve non linearly separable problems



- Unless the input is transformed in a better representation

Adding one more layer

- Single hidden layer neural network
 - 2-layer neural network: ignoring input units






Figure : Two different visualizations of a 2-layer neural network. In this example: 3 input units, 4 hidden units and 2 output units

- Q: What if using linear activation in hidden layer?


Capacity of neural network

- Single hidden layer neural network
 - Partition the input space into regions




Capacity of neural network

- Single hidden layer neural network
 - Form a stump/delta function



Capacity of neural network

- Single hidden layer neural network




Multi-layer perceptron

■ Boolean case

- Multilayer perceptrons (MLPs) can compute more complex Boolean functions
- MLPs can compute **any** Boolean function
 - Since they can emulate individual gates
- MLPs are *universal Boolean functions*

$$((A \& \bar{X} \& Z) | (A \& \bar{Y})) \& ((X \& Y) | (\overline{X \& Z}))$$



Capacity of neural network

- Universal approximation

- Theorem (Hornik, 1991)

- A single hidden layer neural network with a linear output unit can approximate any continuous function arbitrarily well, **given enough hidden units**.

- The result applies for sigmoid, tanh and many other hidden layer activation functions

- Caveat: good result but not useful in practice

- How many hidden units?
 - How to find the parameters by a learning algorithm?

General neural network

- Multi-layer neural network





Figure : A 3-layer neural net with 3 input units, 4 hidden units in the first and second hidden layer and 1 output unit

- Naming conventions; a N -layer neural network:
 - ▶ $N - 1$ layers of hidden units
 - ▶ One output layer

Multilayer networks



network: *sequence of parallelized weighted sums and non-linearities*

DEFINE $\mathbf{x}^{(0)} \equiv \mathbf{x}$, $\mathbf{x}^{(1)} \equiv \mathbf{h}$, ETC.

1st layer

$$\mathbf{s}^{(1)} = \mathbf{W}^{(1)} \mathbf{x}^{(0)}$$

$$\mathbf{x}^{(1)} = \sigma(\mathbf{s}^{(1)})$$


2nd layer

$$\mathbf{s}^{(2)} = \mathbf{W}^{(2)} \mathbf{x}^{(1)}$$

$$\mathbf{x}^{(2)} = \sigma(\mathbf{s}^{(2)})$$

\dots

Multilayer networks




network: *sequence of parallelized weighted sums and non-linearities*

$$\begin{array}{c} \text{output} \end{array} = \sigma \left(\dots \sigma \left(\begin{array}{c} \text{2nd weights} \end{array} \sigma \left(\begin{array}{c} \text{1st weights} \end{array} \begin{array}{c} \text{input} \end{array} \right) \right) \dots \right)$$

Why more layers (deeper)?

- A deep architecture can represent certain functions more compactly
 - (Montufar et al., NIPS'14)
 - Functions representable with a **deep rectifier net** can require an exponential number of hidden units with a shallow one.



Why more layers (deeper)?

- A deep architecture can represent certain functions more compactly
 - Example: Boolean functions
 - There are Boolean functions which require an exponential number of hidden units in the single layer case
 - require a **polynomial number of hidden units** if we can adapt the number of layers
 - Example: multivariate polynomials (Rolnick & Tegmark, ICLR'18)
 - Total number of neurons m required to approximate natural classes of multivariate polynomials of n variables
 - grows **only linearly with n** for deep neural networks, but grows exponentially when merely a single hidden layer is allowed.

Other network connectivity


sequential connectivity: *information must flow through the entire sequence to reach the output*



information may not be able to propagate easily


→ *make shorter paths to output*

residual & highway
connections



Deep residual learning for image recognition, He et al., 2016
Highway networks, Srivastava et al., 2015

dense (concatenated)
connections



Densely connected convolutional networks, Huang et al., 2017

Outline

- Multi-layer neural networks
 - Limitations of single layer networks
 - Neural networks with single hidden layer
 - Sequential network architecture and variants
- Inference and learning
 - Forward and Backpropagation
 - Examples: one-layer network
 - General BP algorithm

Acknowledgement: Rich Zemel@UofT & Feifei Li's cs231n notes

Computation in neural network


- We only need to know two algorithms
 - Inference/prediction: simply forward pass
 - Parameter learning: needs backward pass
- Basic fact:
 - A neural network is a function of composed operations

$$f_L(\mathbf{w}_L, f_{L-1}(\mathbf{w}_{L-1}, \dots f_1(\mathbf{w}_1, \mathbf{x}) \dots))$$

- All the f functions are linear + (simple) nonlinear (differentiable a.e.) operators

Inference example: Forward Pass

- What does the network compute?




- Output of the network can be written as:

$$h_j(\mathbf{x}) = f(v_{j0} + \sum_{i=1}^D x_i v_{ji})$$
$$o_k(\mathbf{x}) = g(w_{k0} + \sum_{j=1}^J h_j(\mathbf{x}) w_{kj})$$

(j indexing hidden units, k indexing the output units, D number of inputs)

Forward Pass in Python

- Example code for a forward pass for a 3-layer network in Python:



```
# forward-pass of a 3-layer neural network:  
f = lambda x: 1.0/(1.0 + np.exp(-x)) # activation function (use sigmoid)  
x = np.random.randn(3, 1) # random input vector of three numbers (3x1)  
h1 = f(np.dot(W1, x) + b1) # calculate first hidden layer activations (4x1)  
h2 = f(np.dot(W2, h1) + b2) # calculate second hidden layer activations (4x1)  
out = np.dot(W3, h2) + b3 # output neuron (1x1)
```

- Can be implemented efficiently using matrix operations

Parameter learning: Backward Pass

■ Supervised learning framework

- Find weights:

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \sum_{n=1}^N \operatorname{loss}(\mathbf{o}^{(n)}, \mathbf{t}^{(n)})$$

where $\mathbf{o} = f(\mathbf{x}; \mathbf{w})$ is the output of a neural network

- Define a loss function, eg:

- ▶ Squared loss: $\sum_k \frac{1}{2} (o_k^{(n)} - t_k^{(n)})^2$
- ▶ Cross-entropy loss: $-\sum_k t_k^{(n)} \log o_k^{(n)}$

- Gradient descent:

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \frac{\partial E}{\partial \mathbf{w}^t}$$

where η is the learning rate (and E is error/loss)

Backward pass

■ Backpropagation

- An efficient method for computing gradients in NNs
- A neural network as a function of composed operations

$$f_L(\mathbf{w}_L, f_{L-1}(\mathbf{w}_{L-1}, \dots f_1(\mathbf{w}_1, \mathbf{x}) \dots))$$

and the loss \mathcal{L} is a function of the network output

→ use chain rule to calculate gradients

Review: Chain rule

- Formal definition

For any nested function $y = f(g(x))$

$$\frac{dy}{dx} = \frac{\partial y}{\partial g(x)} \frac{dg(x)}{dx}$$

Check - we can confirm that : $\Delta y = \frac{dy}{dx} \Delta x$


$$z = g(x) \Rightarrow \Delta z = \frac{dg(x)}{dx} \Delta x$$

$$y = f(z) \Rightarrow \Delta y = \frac{dy}{dz} \Delta z = \frac{dy}{dz} \frac{dg(x)}{dx} \Delta x$$





Example: Single Layer Network

- Let's take a single layer network




Example: Single Layer Network

- Let's take a single layer network and draw it a bit differently




Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} =$$

Example: Single Layer Network




- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}}$$

- Error gradient is computable for any continuous activation function $g()$, and any continuous error function


Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \underbrace{\frac{\partial E}{\partial o_k}}_{\delta_k^o} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}}$$


Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}} = \delta_k^o \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}}$$

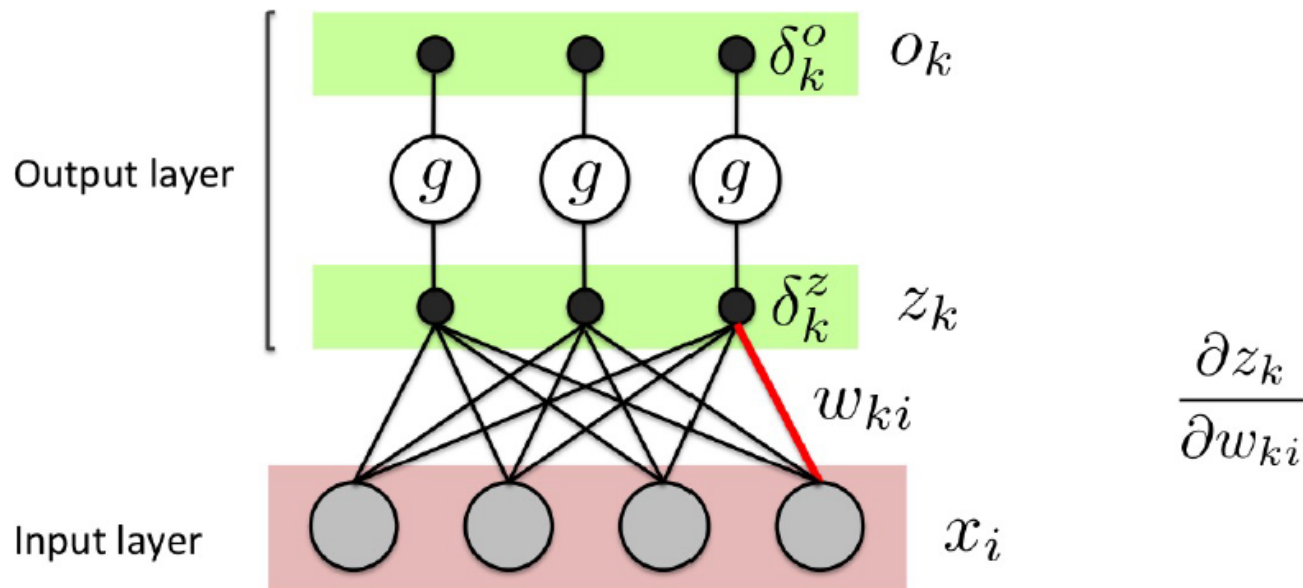
Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}} = \underbrace{\delta_k^o \cdot \frac{\partial o_k}{\partial z_k}}_{\delta_k^z} \frac{\partial z_k}{\partial w_{ki}}$$

Example: Single Layer Network



- Error gradients for single layer network:

$$\frac{\partial E}{\partial w_{ki}} = \frac{\partial E}{\partial o_k} \frac{\partial o_k}{\partial z_k} \frac{\partial z_k}{\partial w_{ki}} = \delta_k^z \frac{\partial z_k}{\partial w_{ki}} = \delta_k^z \cdot x_i$$

Gradient descent iteration


■ Forward pass

1st layer	2nd layer	...	Loss
$\mathbf{s}^{(1)} = \mathbf{W}^{(1)} \tau \mathbf{x}^{(0)}$	$\mathbf{s}^{(2)} = \mathbf{W}^{(2)} \tau \mathbf{x}^{(1)}$		\mathcal{L}
$\mathbf{x}^{(1)} = \sigma(\mathbf{s}^{(1)})$	$\mathbf{x}^{(2)} = \sigma(\mathbf{s}^{(2)})$		

■ Backward pass


calculate $\nabla_{W^{(1)}} \mathcal{L}, \nabla_{W^{(2)}} \mathcal{L}, \dots$ let's start with the final layer: $\nabla_{W^{(L)}} \mathcal{L}$

to determine the chain rule ordering, we'll draw the dependency graph



Gradient descent iteration

■ Backward pass



$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}^{(L)}} \frac{\partial \mathbf{x}^{(L)}}{\partial \mathbf{s}^{(L)}} \frac{\partial \mathbf{s}^{(L)}}{\partial \mathbf{W}^{(L)}}$$

depends on the
form of the loss

derivative of the
non-linearity

$$\frac{\partial}{\partial \mathbf{W}^{(L)}} (\mathbf{W}^{(L)} \mathbf{x}^{(L-1)}) = \mathbf{x}^{(L-1) \top}$$

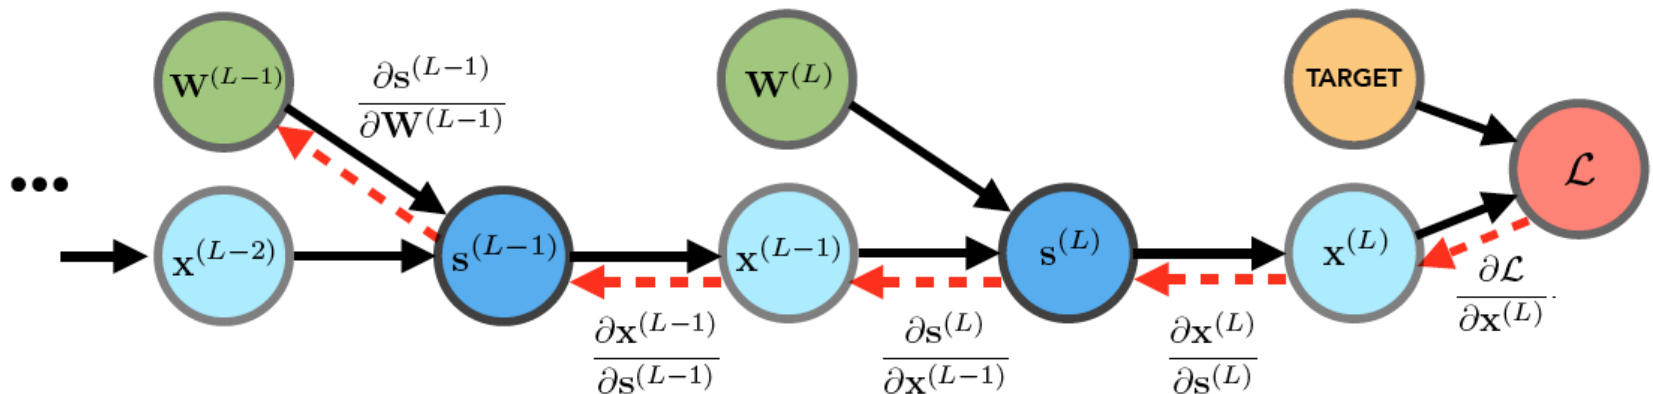
note $\nabla_{\mathbf{W}^{(L)}} \mathcal{L} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L)}}$ is notational convention

Gradient descent iteration

■ Backward pass

now let's go back one more layer...

again we'll draw the dependency graph:



$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(L)}} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}^{(L)}} \frac{\partial \mathbf{x}^{(L)}}{\partial \mathbf{s}^{(L)}} \frac{\partial \mathbf{s}^{(L)}}{\partial \mathbf{x}^{(L-1)}} \frac{\partial \mathbf{x}^{(L-1)}}{\partial \mathbf{s}^{(L-1)}} \frac{\partial \mathbf{s}^{(L-1)}}{\partial \mathbf{W}^{(L-1)}}$$


An implementation perspective

- Example: Univariate logistic least square model

$$s = wx + b$$

$$y = \sigma(s)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$



Univariate chain rule

- A structured way to implement it
 - The goal is to write **a program** that efficiently computes the derivatives

Computing the loss:

$$s = wx + b$$

$$y = \sigma(s)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$

Computing the derivatives:

$$\frac{d\mathcal{L}}{dy} = y - t$$


$$\frac{d\mathcal{L}}{ds} = \frac{d\mathcal{L}}{dy} \sigma'(s)$$

$$\frac{d\mathcal{L}}{dw} = \frac{d\mathcal{L}}{ds} x$$

$$\frac{d\mathcal{L}}{db} = \frac{d\mathcal{L}}{ds}$$

Computation graph

- Represent the computations using a **computation graph**
 - Nodes: inputs & computed quantities
 - Edges: which nodes are computed directly as function of which other nodes



Univariate chain rule

■ A shorthand notation

- Use $\delta_y := d\mathcal{L}/dy$, called the error signal
- Note that the error signals are values computed by the program

Computing the loss:

$$s = wx + b$$

$$y = \sigma(s)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$


Computing the derivatives:

$$\delta_y = y - t$$

$$\delta_s = \delta_y \sigma'(s)$$

$$\delta_w = \delta_s x$$


$$\delta_b = \delta_s$$



Multivariate chain rule

- The computation graph has fan-out > 1

L_2 -Regularized regression



$$z = wx + b$$


$$y = \sigma(z)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$

$$\mathcal{R} = \frac{1}{2}w^2$$

$$\mathcal{L}_{\text{reg}} = \mathcal{L} + \lambda\mathcal{R}$$

Multiclass logistic regression



$$z_\ell = \sum_j w_{\ell j} x_j + b_\ell$$

$$y_k = \frac{e^{z_k}}{\sum_\ell e^{z_\ell}}$$

$$\mathcal{L} = - \sum_k t_k \log y_k$$


Multivariable chain rule

- Recall the distributed chain rule

Mathematical expressions
to be evaluated

$$\frac{df}{dt} = \frac{\partial f}{\partial x} \frac{dx}{dt} + \frac{\partial f}{\partial y} \frac{dy}{dt}$$

Values already computed
by our program



- The shorthand notation:

$$\delta_t = \delta_x \frac{dx}{dt} + \delta_y \frac{dy}{dt}$$

General Backpropagation

- Given a computation graph

Let v_1, \dots, v_N be a **topological ordering** of the computation graph (i.e. parents come before children.)

v_N denotes the variable we're trying to compute derivatives of (e.g. loss)

forward pass


For $i = 1, \dots, N$
Compute v_i as a function of $\text{Pa}(v_i)$

backward pass

$\delta_{v_N} = 1$
For $i = N - 1, \dots, 1$
$$\delta_{v_i} = \sum_{j \in \text{Ch}(v_i)} \delta_{v_j} \frac{\partial v_j}{\partial v_i}$$

General Backpropagation

- Example: univariate logistic least square regression



Forward pass:

$$z = wx + b$$

$$y = \sigma(z)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$

$$\mathcal{R} = \frac{1}{2}w^2$$

$$\mathcal{L}_{\text{reg}} = \mathcal{L} + \lambda \mathcal{R}$$

Backward pass:

$$\delta_{\mathcal{L}_{\text{reg}}} =$$

$$\delta_z =$$

$$\delta_{\mathcal{R}} =$$

$$=$$

$$=$$

$$\delta_w =$$

$$\delta_{\mathcal{L}} =$$

$$=$$

$$=$$

$$\delta y =$$


$$\delta_b =$$

$$=$$

$$=$$

General Backpropagation


- Backprop as message passing:



- Each node receives a set of messages from its children, which are aggregated into its error signal, then it passes messages to its parents
- **Modularity:** each node only has to know how to compute derivatives w.r.t. its arguments – **local computation in the graph**

Patterns in backward flow

- Multiplicative node




Forward:
$$o_i^{(k)} = y_j^{(k-1)} y_l^{(k-1)}$$

$$\frac{\partial L}{\partial y_j^{(k-1)}} = \frac{\partial L}{\partial o_i^{(k)}} \frac{\partial o_i^{(k)}}{\partial y_j^{(k-1)}} = y_l^{(k-1)} \frac{\partial L}{\partial o_i^{(k)}}$$

Patterns in backward flow


■ Max node



- Vector equivalent of subgradient
 - 1 w.r.t. the largest incoming input
 - Incremental changes in this input will change the output
 - 0 for the rest
 - Incremental changes to these inputs will not change the output

Computation cost

- Forward pass: one add-multiply operation per weight
- Backward pass: two add-multiply operations per weight



- For a multilayer network, the cost is linear in the number of layers, quadratic in the number of units per layer

Backpropagation

- Backprop is used to train the majority of neural nets
 - Even generative network learning, or advanced optimization algorithms (second-order) use backprop to compute the update of weights
- However, backprop seems biologically implausible
 - No evidence for biological signals analogous to error derivatives
 - All the existing biologically plausible alternatives learn much more slowly on computers.
 - So how on earth does the brain learn???

Summary

- Multi-layer neural networks
- Inference and learning
 - Forward and Backpropagation
- Next time ...
 - CNN